

10703 Recitation 3

Actor-Critic Methods and Value-Based Methods

Gene Yang

Recap

Definition: The state-value function $v_{\pi}(s)$ of an MDP is the expected return starting from state s , and then following policy π :

- $v_{\pi}(s) = \mathbb{E}[G_t | S_t = s]$

The action-value function $q_{\pi}(s, a)$ is the expected return starting from state s , taking action a , and then following policy:

- $q_{\pi}(s, a) = \mathbb{E}[G_t | S_t = s, A_t = a]$

Learning V with DP

Iterative Policy Evaluation, for estimating $V \approx v_\pi$

Input π , the policy to be evaluated

Algorithm parameter: a small threshold $\theta > 0$ determining accuracy of estimation

Initialize $V(s)$ arbitrarily, for $s \in \mathcal{S}$, and $V(\text{terminal})$ to 0

Loop:

$\Delta \leftarrow 0$

Loop for each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$

Improving π with DP

Policy Improvement

$policy_stable \leftarrow true$

For each $s \in \mathcal{S}$:

$old_action \leftarrow \pi(s)$

$\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s', r} p(s', r | s, a) [r + \gamma V(s')]$

If $old_action \neq \pi(s)$, then $policy_stable \leftarrow false$

If $policy_stable$, then stop and return $V \approx v_*$ and $\pi \approx \pi_*$; else go to 2

Policy Iteration

Policy Iteration (using iterative policy evaluation) for estimating $\pi \approx \pi_*$

1. Initialization

$V(s) \in \mathbb{R}$ and $\pi(s) \in \mathcal{A}(s)$ arbitrarily for all $s \in \mathcal{S}$; $V(\text{terminal}) \doteq 0$

2. Policy Evaluation

Loop:

$\Delta \leftarrow 0$

Loop for each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_{s',r} p(s',r|s,\pi(s)) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$ (a small positive number determining the accuracy of estimation)

3. Policy Improvement

policy-stable \leftarrow true

For each $s \in \mathcal{S}$:

old-action $\leftarrow \pi(s)$

$\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

If *old-action* $\neq \pi(s)$, then *policy-stable* \leftarrow false

If *policy-stable*, then stop and return $V \approx v_*$ and $\pi \approx \pi_*$; else go to 2

Value Iteration

Value Iteration, for estimating $\pi \approx \pi_*$

Algorithm parameter: a small threshold $\theta > 0$ determining accuracy of estimation
Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop:

```
|  $\Delta \leftarrow 0$   
|   Loop for each  $s \in \mathcal{S}$ :  
|      $v \leftarrow V(s)$   
|      $V(s) \leftarrow \max_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$   
|      $\Delta \leftarrow \max(\Delta, |v - V(s)|)$   
until  $\Delta < \theta$ 
```

Output a deterministic policy, $\pi \approx \pi_*$, such that

$$\pi(s) = \operatorname{argmax}_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$$

Q-Learning

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

Initialize S

Choose A from S using policy derived from Q (e.g., ϵ -greedy)

Repeat (for each step of episode):

Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$$

until S is terminal

SARSA

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

Initialize S

Choose A from S using policy derived from Q (e.g., ϵ -greedy)

Repeat (for each step of episode):

Take action A , observe R, S'

Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

until S is terminal

on-policy

Recap

Definition: The state-value function $v_{\pi}(s)$ of an MDP is the expected return starting from state s , and then following policy π :

Deep RL

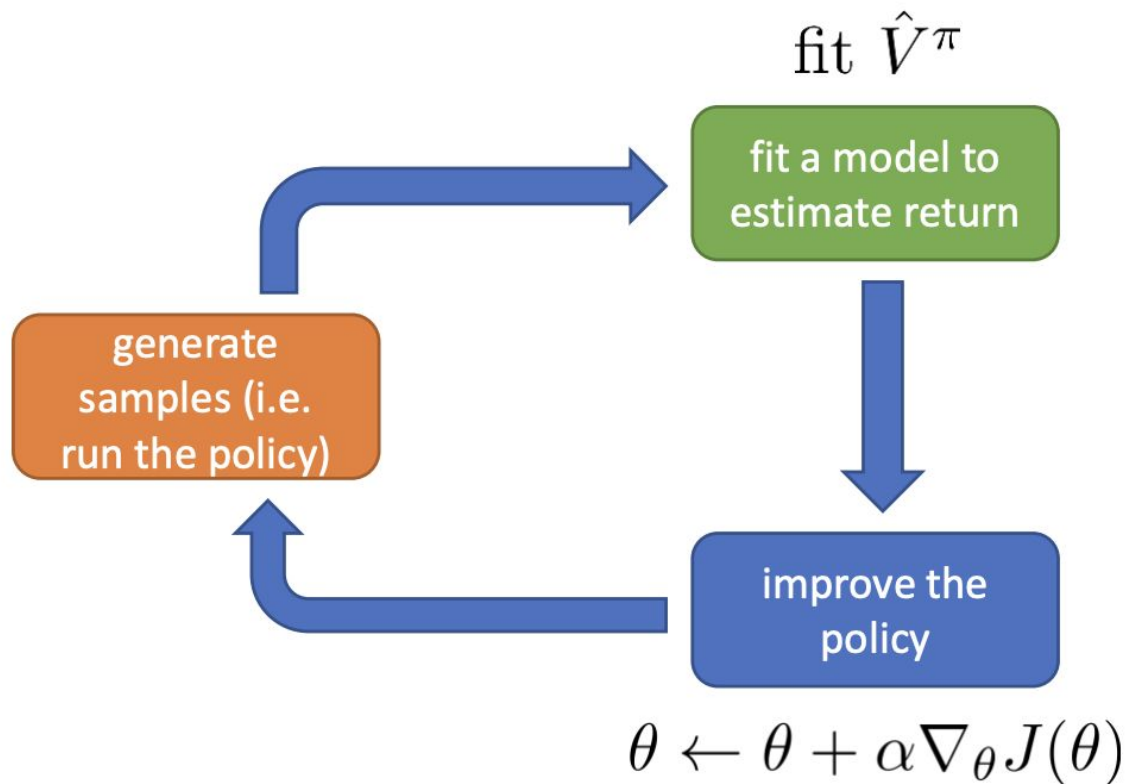
$$V_{\pi}(s; \Phi_1) \text{ or } V_{\Phi_1}^{\pi}(s) \quad \bullet \quad v_{\pi}(s) = \mathbb{E}[G_t | S_t = s]$$

The action-value function $q_{\pi}(s, a)$ is the expected return starting from state s , taking action a , and then following policy:

Deep RL

$$Q_{\pi}(s, a; \Phi_2) \text{ or } Q_{\Phi_2}^{\pi}(s, a) \quad \bullet \quad q_{\pi}(s, a) = \mathbb{E}[G_t | S_t = s, A_t = a]$$

Actor-Critic



Advantage Actor-Critic

0. Initialize policy parameters θ and critic parameters ϕ .

1. Sample trajectories $\{\tau_i = \{s_t^i, a_t^i\}_{t=0}^T\}$ by deploying the current policy $\pi_\theta(a_t | s_t)$.

2. Fit value function $V_\phi^\pi(s)$ by MC or TD estimation (update ϕ)

3. Compute action advantages $A^\pi(s_t^i, a_t^i) = R(s_t^i, a_t^i) + \gamma V_\phi^\pi(s_{t+1}^i) - V_\phi^\pi(s_t^i)$

4. $\nabla_\theta U(\theta) \approx \hat{g} = \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_t^i | s_t^i) A^\pi(s_t^i, a_t^i)$

5. $\theta \leftarrow \theta + \alpha \nabla_\theta U(\theta)$

Advantage Actor-Critic

0. Initialize policy parameters θ and critic parameters ϕ .

1. Sample trajectories $\{\tau_i = \{s_t^i, a_t^i\}_{t=0}^T\}$ by deploying the current policy $\pi_\theta(a_t | s_t)$.

2. Fit value function $V_\phi^\pi(s)$ by MC or TD estimation (update ϕ)

3. Compute action advantages $A^\pi(s_t^i, a_t^i) = R(s_t^i, a_t^i) + \gamma V_\phi^\pi(s_{t+1}^i) - V_\phi^\pi(s_t^i)$

4. $\nabla_\theta U(\theta) \approx \hat{g} = \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_t^i | s_t^i) A^\pi(s_t^i, a_t^i)$ biased, lower variance

5. $\theta \leftarrow \theta + \alpha \nabla_\theta U(\theta)$

DQN

Algorithm 1 - Deep Q-learning with Experience Replay


- 1: Initialize replay memory \mathcal{D} to capacity 50,000 following a random policy.
 - 2: Initialize action-value function Q with random weights θ
 - 3: **for** episode = 1, M **do**
 - 4: **for** $t = 1, T$ **do**
 - 5: With probability ϵ select a random action a_t
 - 6: otherwise select $a_t = \arg \max_a Q(\phi(s_t), a; \theta)$
 - 7: Execute action a_t and observe reward r_t and state s_{t+1}
 - 8: Store (s_t, a_t, r_t, s_{t+1}) in \mathcal{D} .
 - 9: Sample random minibatch of (s_i, a_i, r_i, s_{i+1}) of size N from \mathcal{D}
 - 10: Set $y_j = \begin{cases} r_j & \text{for terminal } s_{j+1} \\ r_j + \gamma \max_{a'} Q(s_{j+1}, a'; \theta) & \text{for non-terminal } s_{j+1} \end{cases}$
 - 11: Perform a gradient descent step on $(y_j - Q(s_j, a_j; \theta))^2$ using Adam
 - 12: **end for**
 - 13: **end for**
-

DQN

Algorithm 1 - Deep Q-learning with Experience Replay

- 1: Initialize replay memory \mathcal{D} to capacity 50,000 following a random policy.
 - 2: Initialize action-value function Q with random weights θ
 - 3: **for** episode = 1, M **do**
 - 4: **for** $t = 1, T$ **do**
 - 5: With probability ϵ select a random action a_t
 - 6: otherwise select $a_t = \arg \max_a Q(\phi(s_t), a; \theta)$
 - 7: Execute action a_t and observe reward r_t and state s_{t+1}
 - 8: Store (s_t, a_t, r_t, s_{t+1}) in \mathcal{D} .
 - 9: decorrelate Sample random minibatch of (s_i, a_i, r_i, s_{i+1}) of size N from \mathcal{D}
 - 10: Set $y_j = \begin{cases} r_j & \text{for terminal } s_{j+1} \\ r_j + \gamma \max_{a'} Q(s_{j+1}, a'; \theta) & \text{for non-terminal } s_{j+1} \end{cases}$
 - 11: Perform a gradient descent step on $(y_j - Q(s_j, a_j; \theta))^2$ using Adam
 - 12: **end for**
 - 13: **end for**
-

- REINFORCE
- Q Learning
- A2C
- A3C
- SARSA
- DQN





on-policy

REINFORCE

A2C

A3C

SARSA

online

???

DQN

Q Learning